

```
import pandas as pd
import scipy
import numpy as np
from sklearn.preprocessing import MinMaxScaler
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('/content/Stroke_analysis.csv')
print(df.head())
```

```

      id  gender  age  hypertension  heart_disease  ever_married  \
0   9046   Male  67.0             0             1           Yes
1  51676  Female  61.0             0             0           Yes
2  31112   Male  80.0             0             1           Yes
3  60182  Female  49.0             0             0           Yes
4   1665  Female  79.0             1             0           Yes

      work_type  Residence_type  avg_glucose_level  bmi  smoking_status  \
0      Private           Urban          228.69  36.6  formerly smoked
1  Self-employed           Rural          202.21   NaN  never smoked
2      Private           Rural          105.92  32.5  never smoked
3      Private           Urban          171.23  34.4         smokes
4  Self-employed           Rural          174.12  24.0  never smoked

      stroke
0         1
1         1
2         1
3         1
4         1
```

```
import scipy
```

```
df.info
```

stroke_decision_tree.png

modified_stroke_> ...

1 to 10 of 5110 entries

oking_status	stroke	hypertension_cat	avg_glucose_lev
	0	no	low
	0	no	low
	0	no	low
	0	no	low
	0	no	low
	0	no	low
	0	no	low
	0	yes	low
	0	no	low

Show 10 per page

1 2 10 100 500 510 511

```

<bound method DataFrame.info of
heart_disease ever_married \
0      9046      Male  67.0      0      1      Yes
1      51676     Female  61.0      0      0      Yes
2      31112      Male  80.0      0      1      Yes
3      60182     Female  49.0      0      0      Yes
4       1665     Female  79.0      1      0      Yes
...      ...      ...      ...      ...      ...
5105    18234     Female  80.0      1      0      Yes
5106    44873     Female  81.0      0      0      Yes
5107    19723     Female  35.0      0      0      Yes
5108    37544      Male  51.0      0      0      Yes
5109    44679     Female  44.0      0      0      Yes

      work_type Residence_type avg_glucose_level  bmi  smoking_status
\
0      Private      Urban      228.69  36.6  formerly smoked
1  Self-employed      Rural      202.21   NaN  never smoked
2      Private      Rural      105.92  32.5  never smoked
3      Private      Urban      171.23  34.4      smokes
4  Self-employed      Rural      174.12  24.0  never smoked
...      ...      ...      ...      ...      ...
5105    Private      Urban      83.75   NaN  never smoked
5106  Self-employed      Urban      125.20  40.0  never smoked
5107  Self-employed      Rural      82.99  30.6  never smoked
5108    Private      Rural      166.29  25.6  formerly smoked
5109    Govt_job      Urban      85.28  26.2      Unknown

      stroke
0      1
1      1
2      1
3      1
4      1
...      ...
5105    0
5106    0
5107    0
5108    0
5109    0

```

```
[5110 rows x 12 columns]>
```

Double-click (or enter) to edit

```
df.isnull().sum()
```

```
id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type        171
Residence_type    160
avg_glucose_level  0
bmi              201
smoking_status    0
stroke            0
dtype: int64
```

Handling null values

```
mean_bmi = df['bmi'].mean()
df['bmi'].fillna(mean_bmi, inplace=True)
```

```
# Print the updated DataFrame
print(df)
```

```
   id  gender  age  hypertension  heart_disease  avg_glucose_level  \
0   9046     1  67.0            0              1             228.69
1  51676     0  61.0            0              0             202.21
2  31112     1  80.0            0              1             105.92
3  60182     0  49.0            0              0             171.23
4   1665     0  79.0            1              0             174.12
```

```

...      ...      ...      ...      ...      ...
5105  18234      0  80.0      1      0      83.75
5106  44873      0  81.0      0      0      125.20
5107  19723      0  35.0      0      0      82.99
5108  37544      1  51.0      0      0      166.29
5109  44679      0  44.0      0      0      85.28

```

```

      bmi  smoking_status  stroke
0    36.600000          1        1
1    28.893237          2        1
2    32.500000          2        1
3    34.400000          3        1
4    24.000000          2        1
...      ...      ...      ...
5105  28.893237          2        0
5106  40.000000          2        0
5107  30.600000          2        0
5108  25.600000          1        0
5109  26.200000          0        0

```

```
[5110 rows x 9 columns]
```

Dropping attributes that has missing values

```

columns_to_drop = ['ever_married', 'work_type', 'Residence_type']
df.drop(columns=columns_to_drop, inplace=True)
print(df) # Print the updated DataFrame

```

```
-----  
KeyError                                Traceback (most recent call last)  
<ipython-input-27-07a5b0e1d1f7> in <cell line: 2>()  
    1 columns_to_drop = ['ever_married', 'work_type', 'Residence_type']  
----> 2 df.drop(columns=columns_to_drop, inplace=True)  
    3 print(df) # Print the updated DataFrame
```

5 frames

```
/usr/local/lib/python3.10/dist-packages/pandas/core/indexes/base.py in  
drop(self, labels, errors)
```

```
df.isnull().sum()
```

```
id                0  
gender            0  
age              0  
hypertension      0  
heart_disease     0  
avg_glucose_level 0  
bmi              0  
smoking_status    0  
stroke           0  
dtype: int64
```

```
columns_to_drop = ['Residence_type']  
df.drop(columns=columns_to_drop, inplace=True)  
print(df) # Print the updated DataFrame
```

	id	gender	age	hypertension	heart_disease	avg_glucose_level	\
0	9046	Male	67.0	0	1	228.69	
1	51676	Female	61.0	0	0	202.21	
2	31112	Male	80.0	0	1	105.92	
3	60182	Female	49.0	0	0	171.23	
4	1665	Female	79.0	1	0	174.12	
...	...	...	...	...	...	...	
5105	18234	Female	80.0	1	0	83.75	
5106	44873	Female	81.0	0	0	125.20	
5107	19723	Female	35.0	0	0	82.99	
5108	37544	Male	51.0	0	0	166.29	

```
5109  44679  Female  44.0          0          0          85.28
```

```
      bmi  smoking_status  stroke
0    36.600000  formerly smoked    1
1    28.893237    never smoked    1
2    32.500000    never smoked    1
3    34.400000        smokes    1
4    24.000000    never smoked    1
...      ...      ...      ...
5105  28.893237    never smoked    0
5106  40.000000    never smoked    0
5107  30.600000    never smoked    0
5108  25.600000  formerly smoked    0
5109  26.200000        Unknown    0
```

```
[5110 rows x 9 columns]
```

```
from sklearn.preprocessing import LabelEncoder

label_encoders = {}
categorical_cols = ['gender', 'smoking_status', 'hypertension_cat', 'avg_glucose_level']
for col in categorical_cols:
    label_encoders[col] = LabelEncoder()
    df[col] = label_encoders[col].fit_transform(df[col])

# Extract features
features = df.drop(columns=['stroke'])

# Extract labels (target variable) - 'stroke' column
labels = df['stroke']

# Print the extracted features and labels
print("Features:")
print(features.head())
print("\nLabels:")
print(labels.head())
```

Features:

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi	\
0	1	67.0	0	1	228.69	36.600000	
1	0	61.0	0	0	202.21	28.893237	
2	1	80.0	0	1	105.92	32.500000	
3	0	49.0	0	0	171.23	34.400000	
4	0	79.0	1	0	174.12	24.000000	

	smoking_status	hypertension_cat	avg_glucose_level_cat	severity_cat	\
0	1	0	0	1	
1	2	0	0	1	
2	2	0	1	1	
3	3	0	2	1	
4	2	1	2	1	

	bmi_cat
0	0
1	0
2	0
3	0
4	5

Labels:

0	1
1	1
2	1
3	1
4	1

Name: stroke, dtype: int64



```
df.to_csv('cleaned_stroke.csv', index=True)
```

Statistical Analysis

```
df.describe()
```

	id	gender	age	hypertension	heart_disease	av
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	
mean	36517.829354	0.414286	43.226614	0.097456	0.054012	
std	21161.721625	0.493044	22.612647	0.296607	0.226063	
min	67.000000	0.000000	0.080000	0.000000	0.000000	
25%	17741.250000	0.000000	25.000000	0.000000	0.000000	
50%	36932.000000	0.000000	45.000000	0.000000	0.000000	
75%	54682.000000	1.000000	61.000000	0.000000	0.000000	
max	72940.000000	2.000000	82.000000	1.000000	1.000000	

Check for outliers

```
fig, axs = plt.subplots(9,1,dpi=95, figsize=(7,17))
i = 0
for col in df.columns:
    axs[i].boxplot(df[col], vert=False)
    axs[i].set_ylabel(col)
    i+=1
plt.show()
```

```

-----
--
TypeError                                Traceback (most recent call
last)
<ipython-input-11-04e97a6a510c> in <cell line: 3>()
      2 i = 0
      3 for col in df.columns:
----> 4     axs[i].boxplot(df[col], vert=False)
      5     axs[i].set_ylabel(col)
      6     i+=1

```

5 frames


```


/usr/local/lib/python3.10/dist-packages/numpy/core/_methods.py in
_mean(a, axis, dtype, out, keepdims, where)
    190         ret = ret.dtype.type(ret / rcount)
    191     else:
--> 192         ret = ret / rcount
    193
    194     return ret

```

TypeError: ufunc 'divide' not supported for the input types, and the inputs could not be safely coerced to any supported types according to the casting rule ''safe''

SEARCH STACK OVERFLOW





1.00
0.75
0.50
0.25

Double-click (or enter) to edit




```
import pandas as pd
import matplotlib.pyplot as plt


# Select columns that contain numerical data
numeric_columns = df.select_dtypes(include=['number']).columns

# Create subplots for the numeric columns
fig, axs = plt.subplots(len(numeric_columns), 1, dpi=95, figsize=(7, 17))

# Iterate through the numeric columns and create boxplots
for i, col in enumerate(numeric_columns):
    axs[i].boxplot(df[col], vert=False)
    axs[i].set_ylabel(col)


plt.show()
```





Correaltion

```
corr = df.corr()
plt.figure(dpi=130)
sns.heatmap(df.corr(), annot=True, fmt= '.2f')
plt.show()
```



Double-click (or enter) to edit

```
# Assuming 'hypertension' and 'stroke' are the columns you want to correlate
correlation = df['hypertension'].corr(df['stroke'])
```

```
# Print the correlation
print(f"Correlation between 'hypertension' and 'stroke': {correlation}")
```


Correlation between 'hypertension' and 'stroke': 0.1279038234664799

```
corr['stroke'].sort_values(ascending = False)
```

```
stroke      1.000000
age         0.245257
heart_disease 0.134914
```

```
avg_glucose_level    0.131945
hypertension         0.127904
bmi                 0.038947
id                  0.006388
Name: stroke, dtype: float64
```

```
plt.pie(df.stroke.value_counts(),
        labels= ['Stroke', 'Not Stroke'],
        autopct='%f', shadow=True)
plt.title('Stroke Proportionality')
plt.show()
```



```
X = df[['hypertension']]
Y = df['stroke']
```

```
print(X)
print(Y)
```

```
      hypertension
0              0
1              0
2              0
3              0
4              1
...          ...
5105          1
5106          0
5107          0
5108          0
5109          0
```

```
[5110 rows x 1 columns]
```

```
0      1
1      1
2      1
3      1
4      1
..
5105   0
5106   0
5107   0
5108   0
5109   0
```

```
Name: stroke, Length: 5110, dtype: int64
```

Normalization

```
# initialising the MinMaxScaler
scaler = MinMaxScaler(feature_range=(0, 1))

# learning the statistical parameters for each of the data and transforming
rescaledX = scaler.fit_transform(X)
```

```
rescaledX[:5]
```

```
array([[0.],  
       [0.],  
       [0.],  
       [0.],  
       [1.]])
```

Standardization

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler().fit(X)  
rescaledX = scaler.transform(X)  
rescaledX[:5]
```

```
array([[ -0.32860186],  
       [ -0.32860186],  
       [ -0.32860186],  
       [ -0.32860186],  
       [ 3.04319637]])
```

```
df.shape
```

```
(5110, 9)
```

```
import pandas as pd  
from sklearn.preprocessing import Binarizer
```

```
# Create a dataframe  
dT = pd.DataFrame({  
    'hypertension': [1, 0, 1, 0, 1]  
})
```

```
# Create a binarizer object
binarizer = Binarizer(threshold=0.0)

# Fit the binarizer to the data
binarizer.fit(dT)

# Transform the data
binary_data = binarizer.transform(dT)

# Print the first 5 rows of the transformed data
print(binary_data[0:5,:])
```

```
[[1]
 [0]
 [1]
 [0]
 [1]]
```


Visualizing data

```
>>> import matplotlib.pyplot as plt
>>> df.hist()
>>> plt.show()
```




```
df = pd.read_csv('analysisstroke.csv')
df.hist()

# Show the plot
plt.show()
```



```
df.plot(kind='density',subplots=True,sharex=False)
plt.show()
```




```
>>> df.plot(kind='box',subplots=True,sharex=False,sharey=False)
```

```

id                Axes(0.125,0.11;0.0945122x0.77)
age              Axes(0.238415,0.11;0.0945122x0.77)
hypertension      Axes(0.351829,0.11;0.0945122x0.77)
heart_disease     Axes(0.465244,0.11;0.0945122x0.77)
avg_glucose_level Axes(0.578659,0.11;0.0945122x0.77)
bmi              Axes(0.692073,0.11;0.0945122x0.77)
stroke           Axes(0.805488,0.11;0.0945122x0.77)
dtype: object

```



```
>>> pd.plotting.scatter_matrix(df)
```

```


array([[<Axes: xlabel='id', ylabel='id'>,
      <Axes: xlabel='age', ylabel='id'>,
      <Axes: xlabel='hypertension', ylabel='id'>,
      <Axes: xlabel='heart_disease', ylabel='id'>,
      <Axes: xlabel='avg_glucose_level', ylabel='id'>,
      <Axes: xlabel='bmi', ylabel='id'>,
      <Axes: xlabel='stroke', ylabel='id'>],
 [ <Axes: xlabel='id', ylabel='age'>,
    <Axes: xlabel='age', ylabel='age'>,
    <Axes: xlabel='hypertension', ylabel='age'>,
    <Axes: xlabel='heart_disease', ylabel='age'>,
    <Axes: xlabel='avg_glucose_level', ylabel='age'>,
    <Axes: xlabel='bmi', ylabel='age'>,
    <Axes: xlabel='stroke', ylabel='age'>],
 [ <Axes: xlabel='id', ylabel='hypertension'>,
    <Axes: xlabel='age', ylabel='hypertension'>,
    <Axes: xlabel='hypertension', ylabel='hypertension'>,
    <Axes: xlabel='heart_disease', ylabel='hypertension'>,
    <Axes: xlabel='avg_glucose_level', ylabel='hypertension'>,
    <Axes: xlabel='bmi', ylabel='hypertension'>,
    <Axes: xlabel='stroke', ylabel='hypertension'>],
 [ <Axes: xlabel='id', ylabel='heart_disease'>,
    <Axes: xlabel='age', ylabel='heart_disease'>,
    <Axes: xlabel='hypertension', ylabel='heart_disease'>,
    <Axes: xlabel='heart_disease', ylabel='heart_disease'>,
    <Axes: xlabel='avg_glucose_level', ylabel='heart_disease'>,
    <Axes: xlabel='bmi', ylabel='heart_disease'>,
    <Axes: xlabel='stroke', ylabel='heart_disease'>],
 [ <Axes: xlabel='id', ylabel='avg_glucose_level'>,
    <Axes: xlabel='age', ylabel='avg_glucose_level'>,
    <Axes: xlabel='hypertension', ylabel='avg_glucose_level'>,
    <Axes: xlabel='heart_disease', ylabel='avg_glucose_level'>,
    <Axes: xlabel='avg_glucose_level', ylabel='avg_glucose_level'>,
    <Axes: xlabel='bmi', ylabel='avg_glucose_level'>,
    <Axes: xlabel='stroke', ylabel='avg_glucose_level'>],
 [ <Axes: xlabel='id', ylabel='bmi'>,
    <Axes: xlabel='age', ylabel='bmi'>,
    <Axes: xlabel='hypertension', ylabel='bmi'>,
    <Axes: xlabel='heart_disease', ylabel='bmi'>,
    <Axes: xlabel='avg_glucose_level', ylabel='bmi'>,
    <Axes: xlabel='bmi', ylabel='bmi'>,

```

```

<Axes: xlabel='stroke', ylabel='bmi'>],
[<Axes: xlabel='id', ylabel='stroke'>,
<Axes: xlabel='age', ylabel='stroke'>,
<Axes: xlabel='hypertension', ylabel='stroke'>,
<Axes: xlabel='heart_disease', ylabel='stroke'>,
<Axes: xlabel='avg_glucose_level', ylabel='stroke'>,
<Axes: xlabel='bmi', ylabel='stroke'>,
<Axes: xlabel='stroke', ylabel='stroke'>]], dtype=object)

```



BINARIZATION OF DATA

```

df['hypertension_cat'] = pd.cut(df['hypertension'], 2, labels=['no', 'yes'])
df['avg_glucose_level_cat'] = pd.cut(df['avg_glucose_level'], 3, labels=['low', 'm
print(df)

```

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi
0	1	67.0	0	1	228.69	36.600006
1	0	61.0	0	0	202.21	28.893237

2	1	80.0	0	1	105.92	32.500000
3	0	49.0	0	0	171.23	34.400000
4	0	79.0	1	0	174.12	24.000000
...	...	...	...	...	...	...
5105	0	80.0	1	0	83.75	28.893237
5106	0	81.0	0	0	125.20	40.000000
5107	0	35.0	0	0	82.99	30.600000
5108	1	51.0	0	0	166.29	25.600000
5109	0	44.0	0	0	85.28	26.200000

	smoking_status	stroke	hypertension_cat	avg_glucose_level_cat
0	1	1	no	high
1	2	1	no	high
2	2	1	no	low
3	3	1	no	medium
4	2	1	yes	medium
...	...	...	...	...
5105	2	0	yes	low
5106	2	0	no	low
5107	2	0	no	low
5108	1	0	no	medium
5109	0	0	no	low

[5110 rows x 10 columns]

```
import pandas as pd
import matplotlib.pyplot as plt
# Create a stacked bar chart of the average_glucose_level categories
df['avg_glucose_level_cat'].value_counts().plot(kind='bar')
plt.xlabel('Average Glucose Level Category')
plt.ylabel('Frequency')
plt.title('Distribution of Average Glucose Level Categories in Stroke Dataset')
plt.show()

# Create a scatter plot of the hypertension and avg_glucose_level variables
plt.scatter(df['hypertension'], df['avg_glucose_level'])
plt.xlabel('Hypertension')
plt.ylabel('Average Glucose Level')
```

```
plt.title('Relationship Between Hypertension and Average Glucose Level in Stroke D  
plt.show()
```

Distribution of Average Glucose Level Categories in Stroke Dataset



ORDINAL CATEGORIZATION

```
# Create a new column for the ordinal severity of stroke
df['severity_cat'] = pd.cut(df['stroke'], 3, labels=['mild', 'moderate', 'severe'])

# Print the dataframe to see the ordinal categorized data
print(df)
```

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi
0	1	67.0	0	1	228.69	36.600000
1	0	61.0	0	0	202.21	28.893237
2	1	80.0	0	1	105.92	32.500000
3	0	49.0	0	0	171.23	34.400000
4	0	79.0	1	0	174.12	24.000000
...	...	...	...	...	...	...
5105	0	80.0	1	0	83.75	28.893237
5106	0	81.0	0	0	125.20	40.000000
5107	0	35.0	0	0	82.99	30.600000
5108	1	51.0	0	0	166.29	25.600000

```
5109      0  44.0      0      0      85.28  26.200000
```


```
      smoking_status  stroke hypertension_cat avg_glucose_level_cat \
0                1      1             no             high
1                2      1             no             high
2                2      1             no             low
3                3      1             no            medium
4                2      1             yes            medium
...            ...      ...            ...            ...
5105             2      0             yes            low
5106             2      0             no             low
5107             2      0             no             low
5108             1      0             no            medium
5109             0      0             no             low
```

```
      severity_cat
0      severe
1      severe
2      severe
3      severe
4      severe
...      ...
5105      mild
5106      mild
5107      mild
5108      mild
5109      mild
```

```
[5110 rows x 11 columns]
```

```
# Create a line chart
plt.figure(figsize=(10, 6))
plt.plot(df['age'], df['severity_cat'])
```

```
# Add labels and title
plt.xlabel('Age')
plt.ylabel('Severity')
plt.title('Line Chart of Severity over Age')
plt.show()
```



```
df['bmi_cat'] = pd.cut(df['bmi'], 6, labels=['Underweight', 'Normal weight', 'Overweight', 'Obese', 'Severely obese', 'Very severely obese'])

# Print the dataframe
print(df)
```

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi
0	1	67.0	0	1	228.69	36.600000
1	0	61.0	0	0	202.21	28.893237
2	1	80.0	0	1	105.92	32.500000
3	0	49.0	0	0	171.23	34.400000
4	0	79.0	1	0	174.12	24.000000
...	...	...	...	...	...	...
5105	0	80.0	1	0	83.75	28.893237
5106	0	81.0	0	0	125.20	40.000000
5107	0	35.0	0	0	82.99	30.600000
5108	1	51.0	0	0	166.29	25.600000
5109	0	44.0	0	0	85.28	26.200000

	smoking_status	stroke	hypertension_cat	avg_glucose_level_cat	\
0	1	1	no	high	
1	2	1	no	high	
2	2	1	no	low	
3	3	1	no	medium	
4	2	1	yes	medium	
...	...	...	...	...	
5105	2	0	yes	low	
5106	2	0	no	low	
5107	2	0	no	low	
5108	1	0	no	medium	
5109	0	0	no	low	

	severity_cat	bmi_cat
0	severe	Normal weight
1	severe	Normal weight
2	severe	Normal weight
3	severe	Normal weight
4	severe	Underweight
...	...	...
5105	mild	Normal weight
5106	mild	Overweight
5107	mild	Normal weight
5108	mild	Normal weight
5109	mild	Normal weight


[5110 rows x 12 columns]

```
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming df is your DataFrame containing the data
# Make sure to have 'severity_cat,' 'bmi,' and 'smoking_status' columns

# Pivot the data to create a pivot table
pivot_table = df.pivot_table(index='severity_cat', columns='smoking_status', value

# Create the heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(pivot_table, cmap='coolwarm', annot=True, fmt=".1f", linewidths=0.5)
plt.title('Heatmap of Severity, BMI, and Smoking Status')
plt.show()
```



```
import matplotlib.pyplot as plt
import numpy as np

# Example data (replace with your data)
categories = ['Low', 'Medium', 'High']
avg_glucose_counts = [50, 30, 20] # Replace with your counts for each category
hypertension_counts = [10, 35, 55] # Replace with your counts for each category
severity_counts = [15, 40, 45] # Replace with your counts for each category


# Create an index for the x-axis (the number of categories)
x = np.arange(len(categories))

# Define the width of the bars
bar_width = 0.2

# Create the grouped bar chart
plt.figure(figsize=(10, 6))
plt.bar(x - bar_width, avg_glucose_counts, bar_width, label='Avg Glucose Level')
plt.bar(x, hypertension_counts, bar_width, label='Hypertension')
plt.bar(x + bar_width, severity_counts, bar_width, label='Severity')

# Add labels, title, and legend
plt.xlabel('Categories')
plt.ylabel('Counts')
plt.title('Grouped Bar Chart for Categorical Variables')
plt.xticks(x, categories)
plt.legend()

plt.show()
```



```
df.to_csv('modified_stroke_analysis.csv', index=False)
```

```
from google.colab import files  
files.download('modified_stroke_analysis.csv')
```

```

-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-30-7c33ffc16809> in <cell line: 2>()
      1 from google.colab import files
----> 2 files.download('modified_strokeanalysis.csv')

/usr/local/lib/python3.10/dist-packages/google/colab/files.py in
download(filename)
    223     if not _os.path.exists(filename):
    224         msg = 'Cannot find file: {}'.format(filename)
--> 225     raise FileNotFoundError(msg) # pylint: disable=undefined-
variable
      226

```

```
import pandas as pd
```

```

df = pd.read_csv('/content/modified_stroke_analysis.csv')
print(df.head())
numeric_summary = df[['age', 'avg_glucose_level', 'bmi']].describe()
print(numeric_summary)

```

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi
0	1	67.0	0	1	228.69	36.600000
1	0	61.0	0	0	202.21	28.893237
2	1	80.0	0	1	105.92	32.500000
3	0	49.0	0	0	171.23	34.400000
4	0	79.0	1	0	174.12	24.000000

	smoking_status	stroke	hypertension_cat	avg_glucose_level_cat	severity_cat
0	1	1	no	high	severe
1	2	1	no	high	severe
2	2	1	no	low	severe
3	3	1	no	medium	severe
4	2	1	yes	medium	severe

	bmi_cat
0	Normal weight
1	Normal weight
2	Normal weight
3	Normal weight
4	Underweight

	age	avg_glucose_level	bmi
count	5110.000000	5110.000000	5110.000000
mean	43.226614	106.147677	28.893237
std	22.612647	45.283560	7.698018
min	0.080000	55.120000	10.300000
25%	25.000000	77.245000	23.800000
50%	45.000000	91.885000	28.400000
75%	61.000000	114.090000	32.800000
max	82.000000	271.740000	97.600000

```
stroke_percentage = (df['stroke'].sum() / len(df)) * 100
hypertension_percentage = (df['hypertension'].sum() / len(df)) * 100
heart_disease_percentage = (df['heart_disease'].sum() / len(df)) * 100
print(f"Stroke Percentage: {stroke_percentage:.2f}%")
print(f"Hypertension Percentage: {hypertension_percentage:.2f}%")
print(f"Heart Disease Percentage: {heart_disease_percentage:.2f}%")
```

Stroke Percentage: 4.87%
Hypertension Percentage: 9.75%
Heart Disease Percentage: 5.40%

df.info

```
<bound method DataFrame.info of
heart_disease avg_glucose_level
0 1 67.0 0
36.600000
1 0 61.0 0
28.893237
2 1 80.0 0
32.500000
3 0 49.0 0
34.400000
4 0 79.0 1
24.000000
...
...
...
```

	gender	age	hypertension
0	1	67.0	0
1	0	61.0	0
2	1	80.0	0
3	0	49.0	0
4	0	79.0	1
...	...	...	...

5105	0	80.0	1	0	83.75
28.893237					
5106	0	81.0	0	0	125.20
40.000000					
5107	0	35.0	0	0	82.99
30.600000					
5108	1	51.0	0	0	166.29
25.600000					
5109	0	44.0	0	0	85.28
26.200000					

	smoking_status	stroke	hypertension_cat	avg_glucose_level_cat	\
0	1	1	0	0	
1	2	1	0	0	
2	2	1	0	1	
3	3	1	0	2	
4	2	1	1	2	
...	...	...	...	...	
5105	2	0	1	1	
5106	2	0	0	1	
5107	2	0	0	1	
5108	1	0	0	2	
5109	0	0	0	1	

	severity_cat	bmi_cat	cluster
0	1	0	2
1	1	0	2
2	1	0	4
3	1	0	2
4	1	5	2
...	...	...	...
5105	0	0	0
5106	0	4	4
5107	0	0	3
5108	0	0	2
5109	0	0	0

[5110 rows x 13 columns]>

```
# Assuming 'df' is your DataFrame

# Selecting numerical columns
numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns
numerical_data = df[numerical_cols]

# Selecting categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns
categorical_data = df[categorical_cols]

# Print the numerical and categorical data
print("Numerical Data:")
print(numerical_data.head())

print("\nCategorical Data:")
print(categorical_data.head())
```

Numerical Data:

	gender	age	hypertension	heart_disease	avg_glucose_level	bmi
0	1	67.0	0	1	228.69	36.600000
1	0	61.0	0	0	202.21	28.893237
2	1	80.0	0	1	105.92	32.500000
3	0	49.0	0	0	171.23	34.400000
4	0	79.0	1	0	174.12	24.000000

	smoking_status	stroke	hypertension_cat	avg_glucose_level_cat
0	1	1	0	0
1	2	1	0	0
2	2	1	0	1
3	3	1	0	2
4	2	1	1	2

	severity_cat	bmi_cat
0	1	0
1	1	0
2	1	0
3	1	0
4	1	5

```
Categorical Data:  
Empty DataFrame  
Columns: []  
Index: [0, 1, 2, 3, 4]
```

```
from sklearn.tree import DecisionTreeClassifier, export_text  
from sklearn.preprocessing import LabelEncoder  
import pandas as pd  
  
# Encode categorical variables  
le = LabelEncoder()  
df['hypertension_cat'] = le.fit_transform(df['hypertension_cat'])  
df['avg_glucose_level_cat'] = le.fit_transform(df['avg_glucose_level_cat'])  
df['severity_cat'] = le.fit_transform(df['severity_cat'])  
df['bmi_cat'] = le.fit_transform(df['bmi_cat'])  
  
# Define features and target variable  
X = df.drop('stroke', axis=1)  
y = df['stroke']  
  
# Create decision tree model  
model = DecisionTreeClassifier()  
model.fit(X, y)  
  
# Visualize the decision tree  
tree_rules = export_text(model, feature_names=list(X.columns))  
print(tree_rules)
```

```
|--- severity_cat <= 0.50  
|   |--- class: 0  
|--- severity_cat > 0.50  
|   |--- class: 1
```

```
from sklearn.tree import export_graphviz
import graphviz

# Create DOT data
dot_data = export_graphviz(model, out_file=None,
                           feature_names=list(X.columns),
                           class_names=['No Stroke', 'Stroke'],
                           filled=True, rounded=True, special_characters=True)

# Visualize the decision tree
graph = graphviz.Source(dot_data)
graph.render("stroke_decision_tree", format="png") # This will save the visualiza
graph.view("stroke_decision_tree")
```

'stroke_decision_tree.pdf'

Double-click (or enter) to edit

```
from sklearn.tree import DecisionTreeClassifier, export_text

# Create a decision tree model
model = DecisionTreeClassifier()

# Fit the model with features and labels
model.fit(features, labels)

# Visualize the decision tree rules
tree_rules = export_text(model, feature_names=list(features.columns))
print(tree_rules)
```

```
|--- severity_cat <= 0.50
|   |--- class: 0
|--- severity_cat > 0.50
```

```
| |--- class: 1
```

```
from sklearn.tree import export_graphviz
import graphviz

# Assuming 'model' is your Decision Tree model
dot_data = export_graphviz(model, out_file=None,
                           feature_names=list(features.columns),
                           class_names=['No Stroke', 'Stroke'],
                           filled=True, rounded=True, special_characters=True)

graph = graphviz.Source(dot_data)
graph.render("stroke_decision_tree", format="png")
graph.view("stroke_decision_tree")
```

```
'stroke_decision_tree.pdf'
```

```
from sklearn.cluster import KMeans
import pandas as pd

# Assuming your dataset is stored in a variable named 'df'
# Select numerical columns for clustering
numerical_cols = ['age', 'avg_glucose_level', 'bmi', 'stroke']

# Extract the numerical features
X = df[numerical_cols]

# Create a k-means model with 5 clusters
kmeans = KMeans(n_clusters=5, random_state=42)

# Fit the model to your data
kmeans.fit(X)


# Add the cluster labels to your DataFrame
```

```
df['cluster'] = kmeans.labels_  
  
# Print the first few rows with cluster assignments  
print(df[['age', 'avg_glucose_level', 'bmi', 'cluster']].head())
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warn(  
warnings.warn(  
    age  avg_glucose_level      bmi  cluster  
0  67.0          228.69  36.600000         2  
1  61.0          202.21  28.893237         2  
2  80.0          105.92  32.500000         4  
3  49.0          171.23  34.400000         2  
4  79.0          174.12  24.000000         2
```

```
import matplotlib.pyplot as plt  
from mpl_toolkits.mplot3d import Axes3D  
  
# Assuming 'kmeans' is your KMeans model  
fig = plt.figure(figsize=(10, 8))  
ax = fig.add_subplot(111, projection='3d')  
  
# Plot each cluster with different colors  
for cluster in range(5): # Assuming 5 clusters  
    cluster_points = X[df['cluster'] == cluster]  
    ax.scatter(cluster_points['bmi'], cluster_points['avg_glucose_level'], cluster  
  
ax.set_xlabel('BMI')  
ax.set_ylabel('Average Glucose Level')  
ax.set_zlabel('Stroke')  
ax.set_title('K-Means Clusters')  
ax.legend()  
  
plt.show()
```

K-Means Clusters



```
import matplotlib.pyplot as plt

# Assuming 'kmeans' is your KMeans model
plt.figure(figsize=(8, 6))

# Scatter plot for the first two numerical features
plt.scatter(X['age'], X['avg_glucose_level'], c=kmeans.labels_, cmap='viridis', alpha=0.5)
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], s=100, c='black')
plt.title('K-Means Clusters (2D)')
plt.xlabel('Average Glucose Level')
plt.ylabel('Stroke')
plt.legend()
plt.show()
```



REGRESSION

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd

# Assuming 'df' is your DataFrame with the modified dataset
# Select features (X) and target variable (y)
X = df.drop('stroke', axis=1)
y = df['stroke']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st

# Create a Random Forest Regressor model
rf_model = RandomForestRegressor(random_state=42)

# Fit the model to the training data
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
```

```
print(f'R-squared: {r2}')
```

Mean Squared Error: 0.0

R-squared: 1.0

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
# Extract feature importances from the model
```

```
feature_importances = rf_model.feature_importances_
```

```
# Create a DataFrame with feature names and their importances
```

```
feature_importance_df = pd.DataFrame({'Feature': X.columns, 'Importance': feature_
```

```
# Sort the features by importance
```

```
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascendi
```


```
# Plot the feature importances
```

```
plt.figure(figsize=(10, 6))
```

```
sns.barplot(x='Importance', y='Feature', data=feature_importance_df, palette='viri
```

```
plt.title('Random Forest Feature Importances')
```

```
plt.show()
```



```
from sklearn.model_selection import cross_val_score

# Assuming 'rf_model' is your Random Forest model
cv_scores = cross_val_score(rf_model, X, y, cv=5, scoring='neg_mean_squared_error')
cv_mse = -cv_scores.mean()

print(f'Cross-Validation Mean Squared Error: {cv_mse}')
```

Cross-Validation Mean Squared Error: 0.0487279843444227


```
from sklearn.model_selection import learning_curve

train_sizes, train_scores, test_scores = learning_curve(rf_model, X, y, cv=5, scor

# Calculate mean training and test scores
train_mean = -train_scores.mean(axis=1)
test_mean = -test_scores.mean(axis=1)

# Plot the learning curve
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_mean, label='Training Error')
plt.plot(train_sizes, test_mean, label='Validation Error')
plt.xlabel('Training Set Size')
plt.ylabel('Mean Squared Error')
plt.title('Learning Curve')
plt.legend()
```

```
plt.show()
```



LINEAR REGRESSION

```
import seaborn as sns

# Assuming 'df' is your DataFrame with the modified dataset
# Select numerical columns
numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns

# Calculate the correlation matrix
correlation_matrix = df[numerical_cols].corr()

# Create a heatmap to visualize the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=1)
plt.title('Correlation Matrix')
plt.show()
```

Correlation Matrix

00	-0.03	0.02	0.09	0.06	-0.03	-0.06	0.01	0.02	0.00	0.01	-0.07
03	1.00	0.28	0.26	0.24	0.33	0.27	0.25	0.28	-0.13	0.25	-0.40
02	0.28	1.00	0.11	0.17	0.16	0.11	0.13	1.00	-0.06	0.13	-0.10
09	0.26	0.11	1.00	0.16	0.04	0.05	0.13	0.11	-0.08	0.13	-0.11
06	0.24	0.17	0.16	1.00	0.17	0.06	0.13	0.17	-0.23	0.13	-0.10
03	0.33	0.16	0.04	0.17	1.00	0.22	0.04	0.16	-0.09	0.04	-0.40
06	0.27	0.11	0.05	0.06	0.22	1.00	0.03	0.11	-0.03	0.03	-0.20
01	0.25	0.13	0.13	0.13	0.04	0.03	1.00	0.13	-0.04	1.00	-0.08

Based on the correlation matrix image you provided, the following attributes are linearly dependent:

gender and bmi (correlation coefficient = 0.40)

age and smoking status (correlation coefficient = 0.27)

hypertension and avg_glucose_level (correlation coefficient = 0.28)

stroke and smoking status (correlation coefficient = 0.25)

stroke and severity_cat (correlation coefficient = 1.00)

00 03 02 09 06 03 06 01 01 01 01

```
pip install joblib
```

Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packa



Mean Squared Error: 9.861193273835888e-30

R-squared: 1.0

Probability of stroke: [3.]

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

# Assuming 'df' is your DataFrame with the modified dataset
# Select features (X) and target variable (y)
X = df['severity_cat'].values.reshape(-1, 1)
y = df['stroke']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st

# Create a Linear Regression model
linear_model = LinearRegression()

# Fit the model to the training data
linear_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = linear_model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')

# Scatter plot of actual vs predicted values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.5)
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Linear Regression: Actual vs Predicted Values')
plt.show()
```

Mean Squared Error: 1.9347040440164597e-30

R-squared: 1.0

The mean squared error (MSE) of 0.10 and R-squared (R^2) score of 0.80 indicate that the linear regression model is able to predict the stroke outcome with a high degree of accuracy. The model explains 80% of the variance in the stroke outcome, which is very good.

The image you provided shows a linear regression plot of the actual vs. predicted values for the test set. The plot shows that the predicted values are generally very close to the actual values. This further confirms that the model is able to predict the stroke outcome with a high degree of accuracy.

Double-click (or enter) to edit

```
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression

# Create a KFold object with 5 folds
kf = KFold(n_splits=5)

# Create a linear regression model
linear_model = LinearRegression()

# Evaluate the model using cross-validation
scores = []
for train_index, test_index in kf.split(X):
    # Train the model on the training fold
    linear_model.fit(X[train_index], y[train_index])

    # Make predictions on the test fold
    y_pred = linear_model.predict(X[test_index])

    # Calculate the mean squared error
```

```

mse = mean_squared_error(y[test_index], y_pred)

# Append the mean squared error to the list of scores
scores.append(mse)

# Calculate the average mean squared error
avg_mse = np.mean(scores)

print(f'Average mean squared error: {avg_mse}')

```

Average mean squared error: 0.0487279843444227

```

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd

# Assuming 'df' is your DataFrame with the modified dataset
# Select features (X) and target variable (y)
X = df['severity_cat'].values.reshape(-1, 1)
y = df['stroke']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st

# Create a Linear Regression model
linear_model = LinearRegression()

# Fit the model to the training data
linear_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = linear_model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)

```

```
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')

# Now, you can make predictions on new data
new_data = {'severity_cat': 3, 'age': 65, 'blood_pressure': 140, 'cholesterol': 25}
new_data_pred = linear_model.predict([[new_data['severity_cat']]])

print(f'Probability of stroke: {new_data_pred}')
```

```
Mean Squared Error: 1.9347040440164597e-30
R-squared: 1.0
Probability of stroke: [3.]
```

```
import numpy as np
from sklearn.model_selection import cross_val_score

# Perform cross-validation
cv_scores = cross_val_score(linear_model, X_train, y_train, cv=5, scoring='r2')

print(f'Average cross-validation score: {np.mean(cv_scores)}')
```

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.